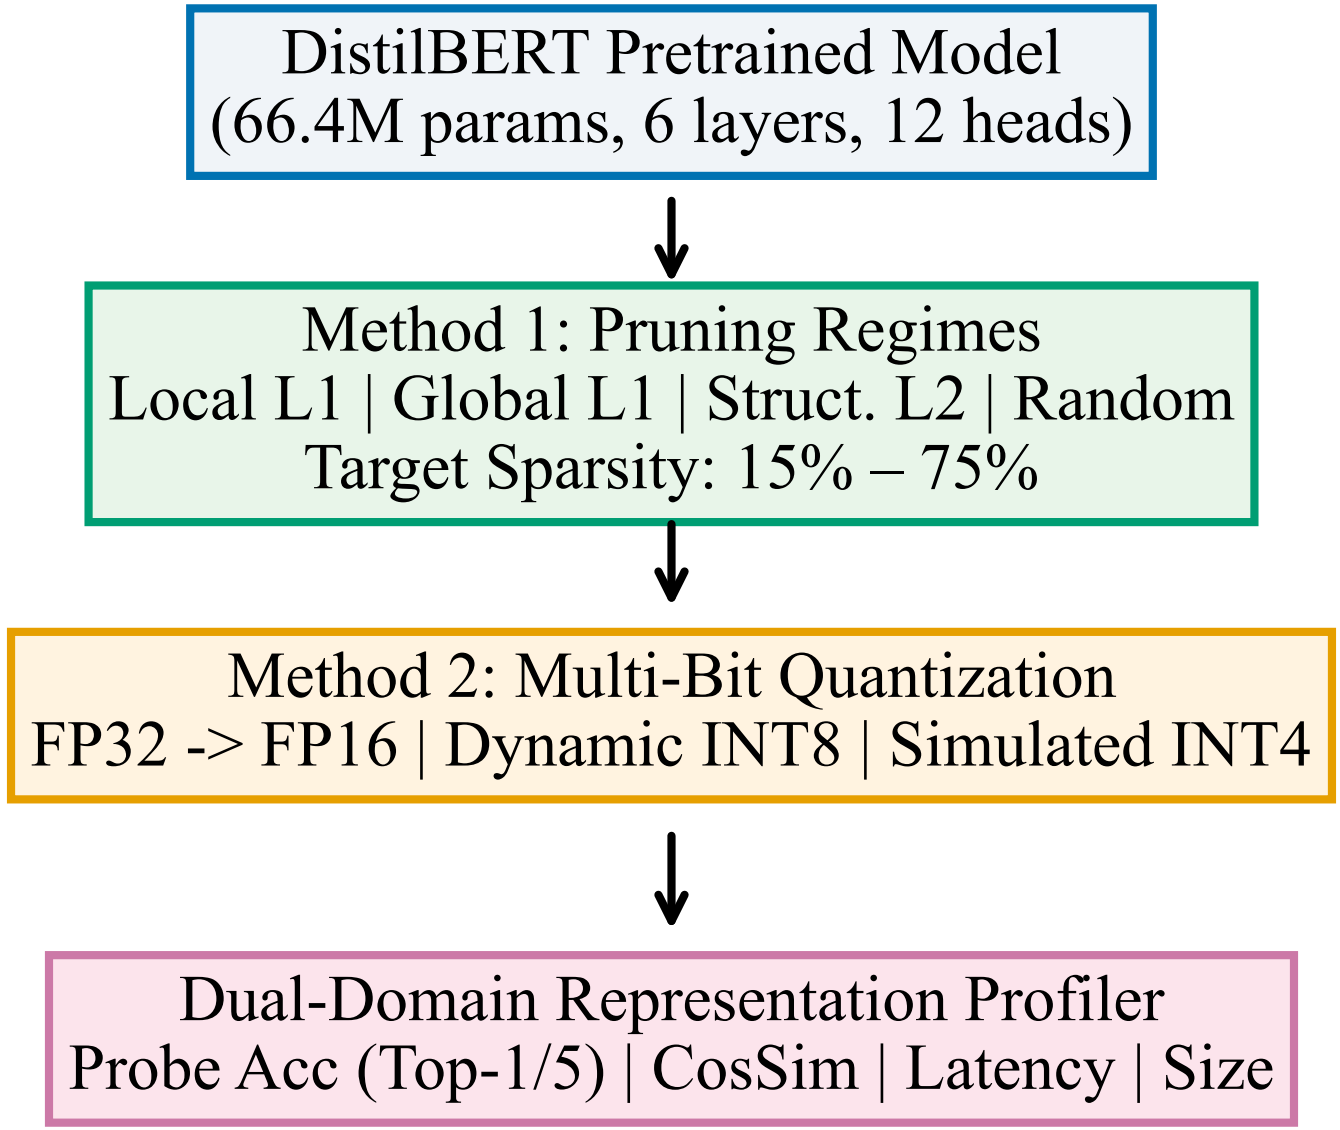
DistilBERT Pretrained Model
(66.4M params, 6 layers, 12 heads)
Method 1: Pruning Regimes
Local L1 | Global L1 | Struct. L2 | Random
Target Sparsity: 15% – 75%
Method 2: Multi-Bit Quantization
FP32 -> FP16 | Dynamic INT8 | Simulated INT4
Dual-Domain Representation Profiler
Probe Acc (Top-1/5) | CosSim | Latency | Size