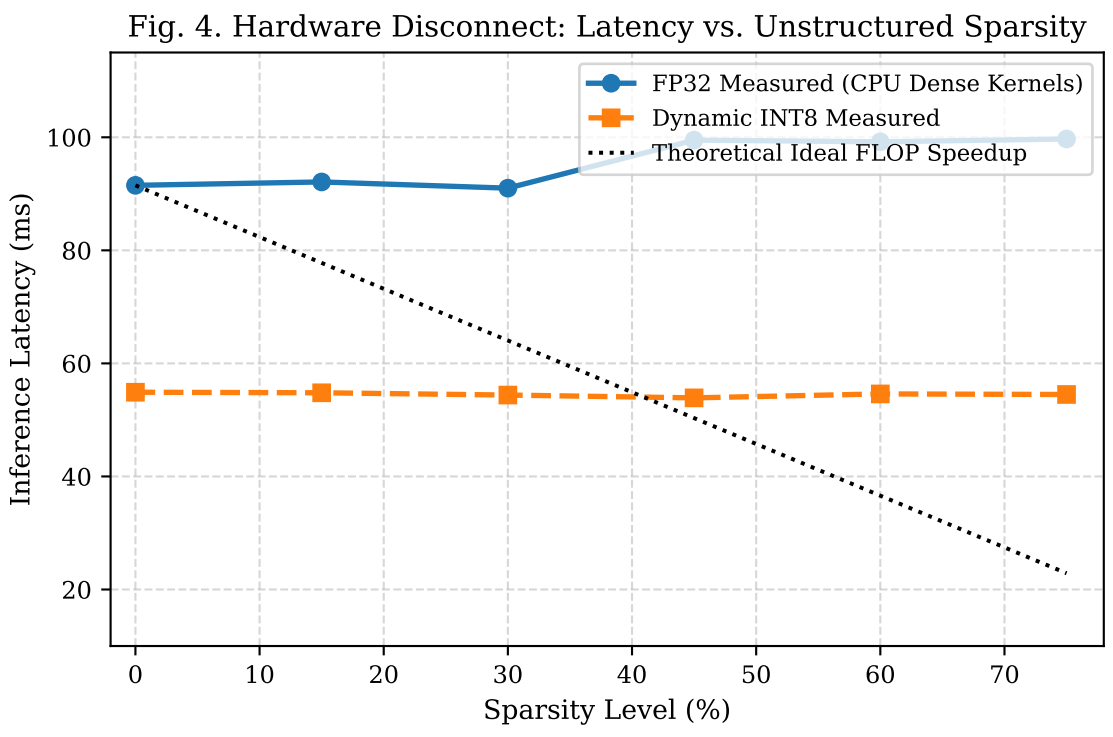
Fig. 4. Hardware Disconnect: Latency vs. Unstructured Sparsity
FP32 Measured (CPU Dense Kernels)
Dynamic INT8 Measured
Theoretical Ideal FLOP Speedup
Inference Latency (ms)
Sparsity Level (%)
100
80
60
40
20
0
10
20
30
40
50
60
70