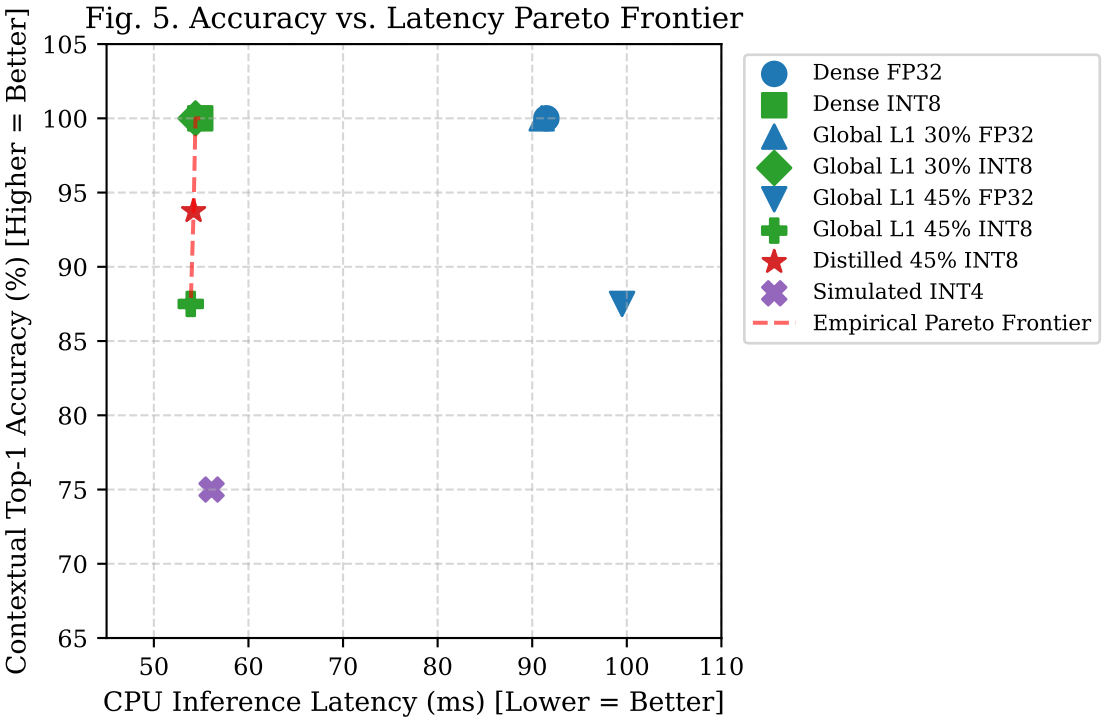
Fig. 5. Accuracy vs. Latency Pareto Frontier
Contextual Top-1 Accuracy (%) [Higher = Better]
CPU Inference Latency (ms) [Lower = Better]
105
100
95
90
85
80
75
70
65
50
60
70
80
90
100
110
Dense FP32
Dense INT8
Global L1 30% FP32
Global L1 30% INT8
Global L1 45% FP32
Global L1 45% INT8
Distilled 45% INT8
Simulated INT4
Empirical Pareto Frontier