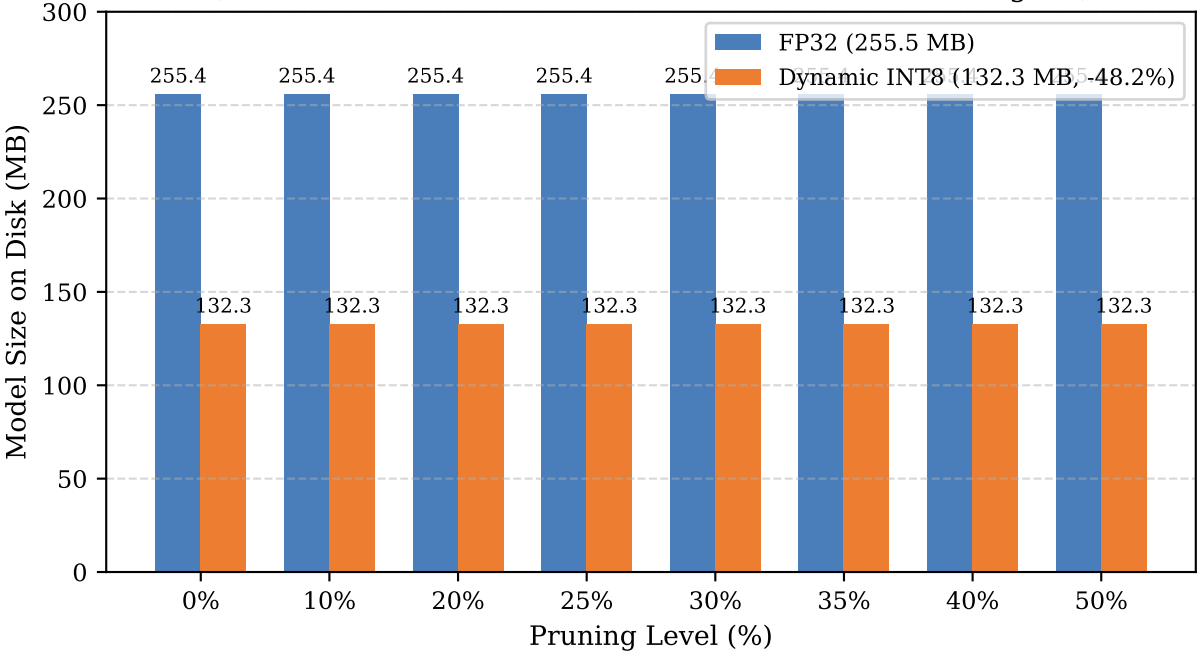
Graph 1 — Model Size vs. Pruning Level
(fine-tuned distilbert-base-uncased-finetuned-sst-2-english)
FP32 (255.5 MB)
Dynamic INT8 (132.3 MB, -48.2%)
Model Size on Disk (MB)
Pruning Level (%)
0%
10%
20%
25%
30%
35%
40%
50%
255.4
132.3
0
50
100
150
200
250
300