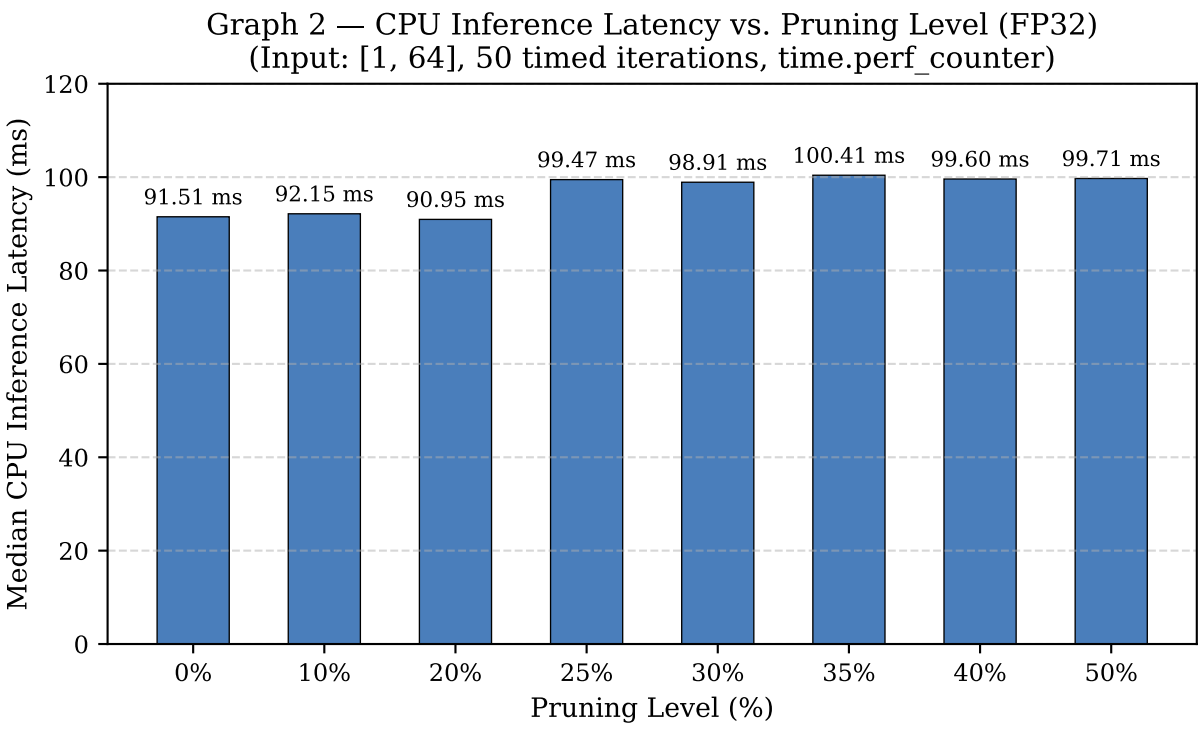
Graph 2 — CPU Inference Latency vs. Pruning Level (FP32)
(Input: [1, 64], 50 timed iterations, time.perf_counter)
Median CPU Inference Latency (ms)
Pruning Level (%)
0
20
40
60
80
100
120
0%
10%
20%
25%
30%
35%
40%
50%
91.51 ms
92.15 ms
90.95 ms
99.47 ms
98.91 ms
100.41 ms
99.60 ms
99.71 ms