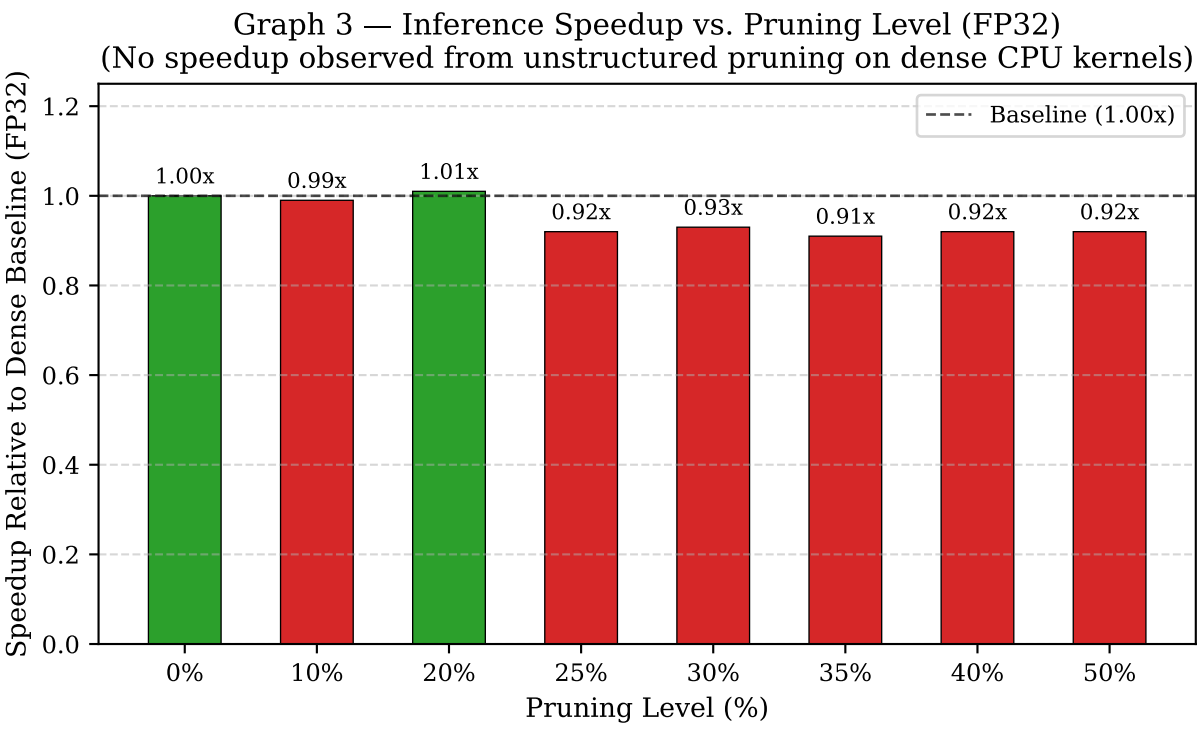
Graph 3 — Inference Speedup vs. Pruning Level (FP32)
(No speedup observed from unstructured pruning on dense CPU kernels)
Speedup Relative to Dense Baseline (FP32)
Pruning Level (%)
Baseline (1.00x)
1.00x
0.99x
1.01x
0.92x
0.93x
0.91x
0.92x
0.92x
0%
10%
20%
25%
30%
35%
40%
50%
0.0
0.2
0.4
0.6
0.8
1.0
1.2